

Prizma-Seq: A Parameter-Free Quadratic Delta-State Sequence Mixer

Aylin

Independent Researcher

Abstract

We present *Prizma-Seq*, a Gated-DeltaNet-family sequence mixer whose only added lever is a *parameter-free quadratic feature map* (QUAD2) that lifts each head’s keys and queries from dimension $d_h=32$ to $d_\phi=256$ using fixed seeded monomial buffers. Because the monomial coefficients are seeded constants rather than learned weights, QUAD2 adds *zero* parameters while making the per-head carried associative state *rectangular* ($d_h \times d_\phi$ rather than $d_h \times d_h$), which preserves exact $O(1)$ -per-step, constant-memory inference. Parameter-matched against a tuned decoder-only Transformer (RMSNorm + SwiGLU + RoPE), Prizma-Seq clears a pre-registered, falsifiable diagnostic bar at small scale: it reaches descriptive parity on MQAR ($D=128$) at $\sim 860\text{K}$ matched parameters, solves $D=128$ at 130K parameters where the matched Transformer fails and the smallest Transformer solver on our coarse grid needs $\geq 461\text{K}$ ($\geq 3.5\times$ parameter-efficiency); it solves induction (0.9995) and selective-copy (0.9991); and it is competitive on a text8 character LM (1.7496 vs. 1.7254 bits-per-character, within a pre-registered $+0.05$ margin — it does *not* beat the Transformer). At inference the carried state is constant in n (0.56 MB / 147,456 floats), analytically $28\text{--}455\times$ fewer floats than a growing KV-cache and a measured constant 17.9 MB peak ($2.8\text{--}31\times$ less than the Transformer); a measured $O(1)$ -latency crossover appears at $n\geq 32\text{k}$ ($2.4\text{--}2.8\times$ faster at $n=65\text{k}$), while below $\sim 16\text{k}$ it is $\sim 1.3\text{--}1.6\times$ *slower*. A causal ablation shows the gain is the quadratic monomials specifically (QUAD2 $\gg$ rand_linear $\approx$ none $\gg$ TF), not “a bigger recurrent state.” We make no per-FLOP, large-scale-LM, or backprop-free-parity claim; Prizma-Seq is a *candidate* efficient-attention replacement in the tested regime, and the QUAD2 kernel itself belongs to the Based/Hedgehog feature-map family — the contribution here is the *rectangular delta-state* framing.

1 Introduction

Self-attention [19]’s score matrix costs $O(n^2)$ compute and an $O(n)$ growing KV-cache at decode time, which dominates the cost of long-context generation. A now-large family of *efficient attention replacements* — linear attention [9], state-space models [8, 7], and delta-rule recurrences [21, 20] — replace the score matrix with a carried associative state of fixed size, trading the $O(n^2)$ matrix for $O(n \cdot d_h^2)$ training and, crucially, $O(d_h^2)$ *per-step, sequence-length-independent* inference. The practical promise is constant-memory, constant-latency decoding at long context.

The cost of this trade is *recall*: a fixed-rank state can store fewer associations than an unbounded KV-cache, and this shows up directly on associative-recall diagnostics such as MQAR [1]. The Based and Hedgehog lines [2, 23] recover much of this recall by passing keys and queries through an expressive (often quadratic) *feature map* before forming the state, enlarging the effective key space.

This paper studies a single, deliberately narrow question: *can a parameter-free quadratic feature map, used to make a Gated-DeltaNet state rectangular, match a tuned Transformer on the field-standard attention-diagnostic suite at matched parameters while preserving $O(1)$ inference?* Our

answer, scoped to small scale and the named tasks, is yes. We call the resulting mixer *Prizma-Seq*. Its lever, QUAD2, lifts each head’s keys/queries from $d_h=32$ to $d_\phi=256$ using *fixed seeded monomial buffers*: the lift is a function with no learnable parameters, so it adds zero to the parameter budget, and it turns the per-head state from a square $d_h \times d_h$ matrix into a *rectangular* $d_h \times d_\phi$ matrix. Because the state size is still independent of n , the constant-memory / $O(1)$ -per-step inference property is preserved exactly.

We test Prizma-Seq against a pre-registered, falsifiable bar (Section 3) and report every number with its honest caveat (Section 7). We emphasize up front what is *not* claimed: Prizma-Seq does not beat the Transformer on the character LM (it is within a margin), the latency advantage is long-context-only, the architecture is $\sim 5\times$ slower to train per step, we make no per-FLOP claim, and we make no large-scale-LM or backprop-free-parity claim. The contribution is the **rectangular delta-state framing** of a known kernel, not a new kernel.

2 Method

2.1 The Gated-DeltaNet backbone

Prizma-Seq replaces self-attention’s token mixer; everything else — the SwiGLU FFN [16], RMSNorm [22], RoPE [17] on the encoder projections, token embeddings, and the tied output head, in the Llama-family style [18] — is byte-identical to our Transformer baseline, so *only the mixer differs*. This makes the comparison a clean attention-replacement test rather than a whole-architecture confound.

Within each head, the mixer carries an associative state S_t updated by a precision-gated *delta rule* [21, 15] with a data-dependent forget gate α_t [20, 7]:

$$S_t = \alpha_t S_{t-1} + \beta_t \varepsilon_t k_t^\top, \quad \varepsilon_t = v_t - S_{t-1} k_t, \quad (1)$$

where k_t are L_2 -normalized keys, v_t values, $\beta_t = \sigma(W_\beta x_t)$ an input-dependent write gate, and ε_t the prediction error (the “surprise”) of the current value under the running state. Reads are recognition-by-reconstruction $o_t = S_{t-1} q_t$, augmented by a small exact local-window head ($w=16$) of the kind used in Based/Griffin [2, 5]. A short causal depthwise convolution (kernel 4), standard in Mamba/Based/DeltaNet, precedes the projections.

Equation (1) is exactly one gradient step on a per-token *free energy* $F_t(S) = \frac{1}{2} \|v_t - S k_t\|^2$, since $\partial F_t / \partial S = -(v_t - S k_t) k_t^\top = -\varepsilon_t k_t^\top$. This predictive-coding reading [14, 6] of the delta write is the project’s framing of the mechanism; the components themselves are borrowed and known-good, and we disclose this explicitly (Section 8). The distinct, separately-tested lever is the feature map below.

2.2 The quad2 parameter-free quadratic feature map

The Based/Hedgehog insight [2, 23] is that passing keys/queries through a quadratic feature map $\phi(\cdot)$ before forming the state recovers much of the recall that a linear state loses. We adopt this kernel family but in a specific, parameter-free form. For a head dimension $d_h=32$, define $\phi : \mathbb{R}^{d_h} \rightarrow \mathbb{R}^{d_\phi}$ with $d_\phi=256$ by *fixed seeded monomials*: the linear terms u_i plus a fixed, seeded selection of second-order cross-products $u_i u_j$,

$$\phi(u) = [u_1, \dots, u_{d_h}, u_{a_1} u_{b_1}, u_{a_2} u_{b_2}, \dots, u_{a_m} u_{b_m}] \in \mathbb{R}^{d_\phi}, \quad (2)$$

where the index pairs $\{(a_r, b_r)\}_{r=1}^m$ are drawn *once* from a fixed seed and stored as constant integer buffers (here $m = 224$ cross-terms plus the 32 linear terms give $d_\phi=256$). Crucially, ϕ has *no learnable parameters*: the monomial structure is a seeded constant, so QUAD2 adds **zero parameters**




Figure 1: The QUAD2 lift makes the per-head delta-state *rectangular* ($d_h \times d_\phi$) at zero added parameters; the per-token write and read stay matrix–vector products over a state whose size $d_h d_\phi$ is independent of the sequence length n , so $O(1)$ inference is preserved.

to the model. Parameter-matching to the Transformer is therefore exact-by-construction — the projection budget is unchanged.

2.3 Rectangular delta-state and why $O(1)$ inference survives

Replacing the keys k_t in Equation (1) with their lifted form $\phi(k_t) \in \mathbb{R}^{d_\phi}$ (and reading with $\phi(q_t)$) turns the carried state from a square $d_h \times d_h$ matrix into a **rectangular** $d_h \times d_\phi$ matrix:

$$S_t = \alpha_t S_{t-1} + \beta_t \varepsilon_t^\phi \phi(k_t)^\top, \quad \varepsilon_t^\phi = v_t - S_{t-1} \phi(k_t), \quad S_t \in \mathbb{R}^{d_h \times d_\phi}, \quad o_t = S_{t-1} \phi(q_t). \quad (3)$$

This is the paper’s contribution: the *rectangular delta-state framing* (Fig. 1). The predictive-coding reading carries over with the lifted key: the write $\beta_t \varepsilon_t^\phi \phi(k_t)^\top$ is one gradient step on $F_t(S) = \frac{1}{2} \|v_t - S \phi(k_t)\|^2$, whose gradient is $\partial F_t / \partial S = -\varepsilon_t^\phi \phi(k_t)^\top$. The recall ceiling is governed by the effective rank of the lifted key space (d_ϕ near-orthogonal directions per head rather than d_h), recovering associative capacity, while the per-token write and read remain a pair of matrix-vector products over a state of *fixed* size $d_h d_\phi$ floats. Because that size does not depend on the sequence length n , the inference profile is unchanged from the linear case: per-step cost is $O(d_h d_\phi)$ and the carried state is constant in n . We verified that the single-step recurrence and the chunk-parallel training form agree ($\|\text{step} - \text{forward}\| < 10^{-6}$: none 3.6×10^{-7} , QUAD2 4.8×10^{-7}), so the cheap inference path is the exact same model as the trained one.

The cost of the lift is FLOPs, not parameters or memory: the $d_\phi=256$ expansion multiplies the delta-state FLOPs and brings Prizma-Seq’s forward cost to $2.14\times$ (as-coded; $1.78\times$ kernel-ideal) the Transformer’s per token at the headline scale (Section 7). We make no per-FLOP claim and disclose this ratio plainly.

3 Experimental setup

Pre-registered falsifiable bar. Before running the headline experiments we committed to a falsifiable bar consisting of the field-standard attention diagnostics plus a structural-advantage and a causal-attribution leg. The pass conditions (e.g. MQAR parity within noise at matched

parameters; induction ≥ 0.98 with the 64–256 gap bin ≥ 0.95 ; selective-copy ≥ 0.97 with a fixed-position control; character LM test $\text{BPC} \leq T + 0.05$; constant-memory inference; and a causal control in which a random *linear* map matches the no-feature-map baseline) were fixed in advance. A leg can fail; the report states whatever the experiment shows.

Baseline and parameter-match auditor. The baseline is a tuned decoder-only Transformer with RMSNorm, SwiGLU, and RoPE, sharing the FFN, norms, embeddings, and tied head with Prizma-Seq so that only the mixer differs. A parameter-match auditor counts both models at each scale; the headline MQAR scale d128L4H4 is matched to **TF 857,216 vs. Prizma-quad2 862,368 params** (+0.6%), and the character-LM scale to TF 3,220,480 vs. Prizma-quad2 3,224,608 (+0.13%). Scale notation is $d\{\text{model}\}L\{\text{layers}\}H\{\text{heads}\}$; throughout, $d_h=32$ and QUAD2 uses $d_\phi=256$.

Protocol and honesty notes. MQAR and induction exhibit a seed-dependent phase transition at this scale *for both* architectures, so we report median, solve-rate (>0.9), and seed range over ≥ 3 seeds rather than a fragile mean. The MQAR/induction/selective-copy/causal legs ran at a single shared $\text{lr}=10^{-3}$ (the Transformer’s tuned value, which is *conservative for* Prizma-Seq, disadvantaging it, not the baseline); only the character LM swept a per-model learning rate. We disclose two reproducibility caveats binding on the diagnostic legs: per-seed weight init is not bit-pinned by the seed in those runs (the comparison is symmetric across architectures, but absolute solve-rates are not bit-reproducible), and the seed counts ($n=3$ diagnostics, $n=2$ character LM) are *descriptive*, not powered equivalence tests. Inference *memory* is analytically counted from closed-form float counts; inference *latency* is measured wall-clock on an A100 (median of 5 reps, 2 warmup). All experiments are small scale on the named tasks — exactly what the bar pre-registered.

Reproducibility and availability. All numbers are produced by reproducible scripts; the raw result JSONs, the experiment configs, and the figure-generation code are committed at <https://github.com/nazmiefearmutcu/Prizma> (see `results/` and `paper/figures/`). Each figure in this paper traces to a named JSON under `results/`; the per-figure provenance map is `paper/figures/figdata/PROVENANCE.md`.

4 Results

Table 1 reports the seven legs of the pre-registered bar with headline numbers. All four field-standard diagnostic legs (MQAR, induction, selective-copy, character LM) pass parameter-matched against the tuned Transformer; the structural-advantage leg passes on memory at every length and on latency only at long context; the causal leg attributes the gain to the quadratic monomials; and length extrapolation is a *relative* win.

MQAR: parity and parameter-efficiency. At $\sim 860\text{K}$ matched parameters (d128L4H4, 3 seeds), all three arms solve $D=128$: Prizma-quad2 seeds [0.9991, 0.9983, 0.9955] (median 0.9983) vs. Transformer [0.9787, 0.9995, 0.9998] (median 0.9995). With $n=3$ at ceiling we cannot resolve a difference — this is descriptive parity (a failure-to-reject), not a tested equivalence. At matched params Prizma-quad2 (0.9983) also exceeds Prizma-none (median 0.9583), and ignites fastest of the three (by step 16–20k, vs. TF 44–80k). At the smaller 130K scale (d64L2H2), Prizma-quad2 solves $D=128$ (3/3, seeds [0.999, 0.997, 0.935], median 0.997) where the param-matched Transformer fails (0/3, 0.016). The smallest Transformer that solves $D=128$ on our coarse 4-point grid

Table 1: Prizma-Seq against the pre-registered, parameter-matched bar (tuned Transformer baseline, A100). Numbers are medians over seeds; see Sections 5–7 for caveats. Char-LM is reported as competitive *within* the +0.05 margin, *not* a win.

Leg	Verdict	Headline
MQAR ($D=128$)	PASS	Parity at 860K matched params (QUAD2 median 0.998, seeds [0.996, 0.998, 0.999] vs. TF 0.9995); solves at 130K (3/3, seeds [0.935, 0.997, 0.999]) where the matched TF fails (0/3) and the smallest grid TF solver needs $\geq 461\text{K} \Rightarrow \geq 3.5\times$ param-efficiency (coarse grid).
Induction	PASS	QUAD2 median 0.9995 (3/3) vs. TF 0.9956. <i>Non-discriminating</i> : Prizma-none also solves (1.0).
Selective-copy	PASS	Selective median 0.9991 ($\geq T-0.02$); fixed-position control isolates content-selectivity (spread 10^{-4}). <i>Non-discriminating</i> : Prizma-none also solves.
Char-LM (text8)	PASS	Prizma 1.7496 vs. TF 1.7254 BPC; margin $-0.024 < +0.05$. Within the bar; does not beat the TF.
Inference	PASS (memory)	Constant 0.56 MB analytic state (28–455 $\times$ fewer floats than KV); measured constant 17.9 MB peak (2.8–31 $\times$ less than TF); $O(1)$ -latency crossover at $n \geq 32\text{k}$ (2.4–2.8 $\times$ faster at $n=65\text{k}$).
Causal ablation	PASS	QUAD2 (3/3, median 0.997) $\gg$ rand_linear (0/3, ≈ 0.59) $\approx$ none (0/3, ≈ 0.52) $\gg$ TF (0/3, 0.016). The gain is the monomials, not “a bigger RNN.”
Length-extrapolation	WIN (relative)	Retention at $8\times$ train length: Prizma 0.398 vs. RoPE-TF 0.041 ($\sim 10\times$). Absolute acc still only ~ 0.40 at $8\times$.

{130, 461, 857, 3300}K is d128L2H4 (461K), giving a $\geq 3.5\times$ parameter-efficiency gap (Fig. 2). We flag that one quad2 seed (0.935) is borderline (clearing the 0.9 bar by 0.035) and that a denser Transformer sweep (200–450K) is needed to pin the true crossover.

Capacity question, resolved honestly. A sufficiently large Transformer *does* solve $D=128$ (d128L2H4/461K: 3/3, 0.9999; d256L4H8/3.3M: 3/3). The tiny-Transformer failure is under-capacity at that size, not “attention cannot recall.” The honest differentiators are parameter-efficiency and $O(1)$ -memory inference — never “attention fails.”

Induction and selective-copy. Both pass parameter-matched (d128L2H4, 0.64% match): induction QUAD2 median 0.9995 (3/3, 0.9995 at the 256-gap bin) vs. TF 0.9956; selective-copy QUAD2 median 0.9991 vs. TF 0.9994, with the fixed-position control spread at 10^{-4} . We disclose that these two legs are *non-discriminating*: Prizma-none also solves them, so the feature-map’s edge is visible in the MQAR-param-efficiency and causal-ablation legs, not here.

Character LM (text8). On text8 [12] (train 10M / val 0.5M / test 0.5M chars, $V=27$, AdamW [11, 10] wd=0.1, val-based early stopping, $n=2$ seeds), Prizma-quad2 reaches median test 1.7496 BPC vs. Transformer 1.7254 BPC — a margin of -0.024 , within the pre-registered +0.05 bar and far below the random baseline of 4.7549. We state plainly: Prizma-Seq is competitive *within* the margin and does **not** beat the Transformer here. Neither arm overfits (final BPCs 1.76 / 1.74). The




Figure 2: MQAR ($D=128$) median accuracy vs. matched parameters (3 seeds; 5 for the 857K TF). Prizma-quad2 solves at 131K where the param-matched Transformer fails (0.016); the smallest Transformer solver on the coarse grid $\{130, 461, 857, 3300\}$ K is d128L2H4 at 461K, giving a $\geq 3.5\times$ parameter-efficiency gap. All points trace to `results/gpu_bench.json`.

Table 2: Measured per-step decode latency (ms) and peak memory (MB) vs. sequence length n on an A100 (small config: d128L4H4, $d_h=32$, $d_\phi=256$). Prizma-Seq latency is flat ($O(1)$); the Transformer grows. The latency crossover is at $n=32k$; below $\sim 16k$ Prizma-Seq is slower (overhead-bound). Peak mem is measured process peak (weights + activations + state), not the 0.56 MB analytic state.

n	4,096	8,192	16,384	32,768	65,536
TF latency (ms)	4.47	4.46	5.39	8.95	16.95
Prizma latency (ms)	6.96	6.95	6.95	6.95	7.06
TF peak mem (MB)	49.5	84.9	152.8	289.1	561.7
Prizma peak mem (MB)	17.9	17.9	17.9	17.9	17.9

Prizma arm took ~ 1889 s vs. the Transformer’s ~ 361 s ($\sim 5.2\times$ slower to train).

5 Inference efficiency

The structural advantage of a fixed-size carried state is constant inference memory. Decoding via `model.step()` (KV-cache for the Transformer, $O(1)$ state for Prizma-Seq) on an A100, the *analytic* carried state is $d_h d_\phi$ floats per head; at the small config it totals 147,456 fp32 floats = 0.56 MB, *constant in n* , against a KV-cache that grows from 4,194,304 floats ($n=4k$) to 67,108,864 ($n=65k$) — an analytic $28.4\times$ to $455\times$ float advantage. This shows up as a *measured* peak process memory (model weights + activations + state) that is a **constant 17.9 MB** for Prizma-Seq at every length, while the Transformer’s measured peak grows from 49.5 MB ($n=4k$) to 561.7 MB ($n=65k$) — a measured peak advantage of $2.8\times$ to $31.4\times$. The larger config shows the same pattern (Prizma 232 MB constant vs. TF 491 MB $\rightarrow$ 4532 MB).

Table 2 and Figure 3 report measured per-step latency. Prizma-Seq’s latency is statistically flat across a $16\times$ span of n (~ 7 ms throughout), while the Transformer grows from 4.47 ms to




Figure 3: Inference profile (A100, small config). Top: per-step decode latency is flat for Prizma-Seq and crosses the Transformer at $n=32k$. Bottom: the carried state is constant in n — 28–455× fewer floats than the KV-cache (right, analytic) and 2.8–31× less measured peak memory (left). Data: `results/gpu_latency.json`.

16.95 ms. The wall-clock *latency* crossover is at $n=32,768$; at $n=65k$ Prizma-Seq is 2.4× faster (small config: 7.06 vs. 16.95 ms) and 2.8× faster on the larger config (13.87 vs. 38.90 ms). **Below $n \lesssim 16k$ Prizma-Seq is ~ 1.3 – $1.6\times$ slower** (it is overhead-bound at short context): the latency win is long-context-only and we disclose it as such. The memory advantage, by contrast, holds at every length.

6 Ablation: is it the quadratic monomials?

To rule out the alternative explanation that the gain is merely “a bigger recurrent state” (a wider d_ϕ), we run a causal control (Fig. 4) at d64L2H2 (130K) on MQAR $D=128$ with three feature-map conditions plus the Transformer, all at fixed $lr=2 \times 10^{-3}$ (3 seeds):

$$\underbrace{\text{QUAD2}}_{3/3, \text{ med } 0.997} \gg \underbrace{\text{rand_linear}}_{0/3, \text{ med } 0.589} \approx \underbrace{\text{none}}_{0/3, \text{ med } 0.518} \gg \underbrace{\text{TF}}_{0/3, 0.016}.$$




Figure 4: Causal control at 130K (d64L2H2, MQAR $D=128$, 3 seeds; error bars span the seed range). Only QUAD2 clears the bar (3/3); `rand_linear` (a $d_\phi=256$ random *linear* lift) behaves like no feature map, and both sit far below QUAD2. The recall gain is the quadratic monomials, not state width. Data: `results/gpu_bench.json`.

The `rand_linear` control is a same-dimension random *linear* map ($\text{rank} \leq d_h$, so it expands d_ϕ to 256 without adding any nonlinear structure). It performs like the no-map baseline, not like QUAD2. This isolates the effect: the recall gain comes from the *quadratic monomials* specifically, not from any $d_\phi=256$ widening of the state. Only QUAD2 solves $D=128$ at 130K parameters.

7 Limitations

We state every limit plainly; the project’s credibility rests on not inflating the claim.

1. **Char-LM is within-margin, not a win.** On text8 Prizma-Seq is competitive *within* the pre-registered ± 0.05 bar (1.7496 vs. 1.7254 BPC) but does **not** beat the Transformer.
2. **Training is $\sim 5\times$ slower per step.** The sequential delta recurrence makes Prizma-Seq $\sim 5\times$ slower to train (char-LM: 1889s vs. 361s per arm).
3. **Latency win is long-context-only.** The measured per-step latency crossover is at $n \geq 32k$; below $\sim 16k$ Prizma-Seq is $\sim 1.3\text{--}1.6\times$ slower. The memory advantage holds at all lengths; the latency advantage does not.
4. **No per-FLOP claim.** Prizma-Seq costs $2.14\times$ (as-coded; $1.78\times$ kernel-ideal) the Transformer’s forward FLOPs per token at the headline scale. The two FLOP-matched Transformer arms (d128L9H4: 1/3; d208L4H4: 0/3, median 0.022) were optimization-confounded (consistent with a learning-rate-transfer artifact, since the same-depth narrower d128L4H4 solves), so they are excluded from the verdict and we make **no per-FLOP claim**. Equal-parameter parity already addresses “is it just compute?”.
5. **Seeds are descriptive, not powered.** $n=3$ (diagnostics) / $n=2$ (char-LM) support failure-to-reject / descriptive parity, not a powered equivalence test. The diagnostic legs ran at a

single shared $\text{lr}=10^{-3}$ (conservative for Prizma-Seq) and their per-seed init is not bit-pinned by the seed.

6. **Length-extrapolation is relative.** Prizma-Seq retains $10\times$ more accuracy than a RoPE Transformer at $8\times$ train length, but its own absolute accuracy at $8\times$ is only ~ 0.40 .
7. **No large-scale-LM parity and no backprop-free parity claim.** We test $\leq 3.2\text{M}$ parameters at character level; large-scale LM parity and backprop-free parity are open frontiers, explicitly not claimed.
8. **The kernel is borrowed.** QUAD2 belongs to the Based/Hedgehog feature-map family; the novelty is the *rectangular delta-state framing*, not the kernel.

8 Related work

Linear and feature-map attention. Linear attention [9, 3] reframes attention as a recurrence with a fixed-size state. Based and Hedgehog [2, 23] restore the recall lost by linearization through expressive feature maps, including quadratic ones; we credit the QUAD2 kernel to this family. Our contribution relative to them is the *rectangular* ($d_h \times d_\phi$) delta-state framing and its parameter-free, fixed-seeded-buffer instantiation.

Delta-rule and gated recurrences. The delta rule connects to fast-weight programmers [15]; DeltaNet [21] parallelizes it over sequence length, and Gated DeltaNet [20] adds a data-dependent gate, as in Mamba [7]. Prizma-Seq’s mixer at the parallelizable level coincides with Gated-DeltaNet plus a short conv and a local window head — all borrowed, known-good components; we make this explicit rather than presenting a new primitive.

State-space models. S4 [8], Mamba [7], and the SSM-attention duality [4] establish constant-state sequence mixing; Griffin [5] combines gated recurrences with local attention, motivating our window head.

Recall diagnostics. MQAR and the Zoology suite [1] are the field-standard probes for the recall–efficiency tradeoff; our pre-registered bar uses them directly.

Predictive coding. The free-energy reading of the delta write draws on predictive coding [14, 6]; the optional local/random-feedback training mode relates to direct feedback alignment [13] and is a secondary axis, not a gate here.

9 Conclusion

Prizma-Seq is a Gated-DeltaNet–family sequence mixer whose only lever is a parameter-free quadratic feature map (QUAD2) that makes the carried delta-state rectangular ($d_h \times d_\phi$) at zero added parameters and with $O(1)$ inference preserved. Parameter-matched against a tuned Transformer it clears a pre-registered diagnostic bar at small scale: descriptive MQAR parity at 860K params, $\geq 3.5\times$ parameter-efficiency at 130K, solved induction and selective-copy, a within-margin character LM, constant-memory inference with a long-context latency crossover, and a causal ablation attributing the gain to the quadratic monomials. It is a *candidate* efficient-attention replacement in the tested regime — not a proven alternative. The open frontiers we do not claim are large-scale LM parity, per-FLOP advantage, and backprop-free parity.

Funding. This work received no external funding.

References

- [1] Simran Arora, Sabri Eyuboglu, Aman Timalsina, Isys Johnson, Michael Poli, James Zou, Atri Rudra, and Christopher Ré. Zoology: Measuring and improving recall in efficient language models. In *International Conference on Learning Representations (ICLR)*, 2024.
- [2] Simran Arora, Sabri Eyuboglu, Michael Zhang, Aman Timalsina, Silas Alberti, Dylan Zinsley, James Zou, Atri Rudra, and Christopher Ré. Simple linear attention language models balance the recall-throughput tradeoff. In *International Conference on Machine Learning (ICML)*, 2024.
- [3] Krzysztof Choromanski, Valerii Likhoshesterov, David Dohan, Xingyou Song, Andreea Gane, Tamas Sarlos, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Lukasz Kaiser, David Belanger, Lucy Colwell, and Adrian Weller. Rethinking attention with performers. In *International Conference on Learning Representations (ICLR)*, 2021.
- [4] Tri Dao and Albert Gu. Transformers are ssms: Generalized models and efficient algorithms through structured state space duality. In *International Conference on Machine Learning (ICML)*, 2024.
- [5] Soham De, Samuel L. Smith, Anushan Fernando, Aleksandar Botev, George Cristian-Muraru, Albert Gu, Ruba Haroun, Leonard Berrada, Yutian Chen, Srivatsan Srinivasan, Guillaume Desjardins, Arnaud Doucet, David Budden, Yee Whye Teh, Razvan Pascanu, Nando De Freitas, and Caglar Gulcehre. Griffin: Mixing gated linear recurrences with local attention for efficient language models. *arXiv preprint arXiv:2402.19427*, 2024.
- [6] Karl Friston. The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 2010.
- [7] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. In *Conference on Language Modeling (COLM)*, 2024.
- [8] Albert Gu, Karan Goel, and Christopher Ré. Efficiently modeling long sequences with structured state spaces. In *International Conference on Learning Representations (ICLR)*, 2022.
- [9] Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and François Fleuret. Transformers are rnns: Fast autoregressive transformers with linear attention. In *International Conference on Machine Learning (ICML)*, 2020.
- [10] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations (ICLR)*, 2015.
- [11] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019.
- [12] Matt Mahoney. Large text compression benchmark. <http://mattmahoney.net/dc/text.html>, 2011.
- [13] Arild Nøkland. Direct feedback alignment provides learning in deep neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2016.

- [14] Rajesh P. N. Rao and Dana H. Ballard. Predictive coding in the visual cortex: A functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience*, 2(1):79–87, 1999.
- [15] Imanol Schlag, Kazuki Irie, and Jürgen Schmidhuber. Linear transformers are secretly fast weight programmers. In *International Conference on Machine Learning (ICML)*, 2021.
- [16] Noam Shazeer. Glu variants improve transformer. *arXiv preprint arXiv:2002.05202*, 2020.
- [17] Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. Roformer: Enhanced transformer with rotary position embedding. *arXiv preprint arXiv:2104.09864*, 2021.
- [18] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- [19] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [20] Songlin Yang, Jan Kautz, and Ali Hatamizadeh. Gated delta networks: Improving mamba2 with delta rule. *arXiv preprint arXiv:2412.06464*, 2024.
- [21] Songlin Yang, Bailin Wang, Yu Zhang, Yikang Shen, and Yoon Kim. Parallelizing linear transformers with the delta rule over sequence length. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [22] Biao Zhang and Rico Sennrich. Root mean square layer normalization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [23] Michael Zhang, Kush Bhatia, Hermann Kumbong, and Christopher Ré. The hedgehog and the porcupine: Expressive linear attentions with softmax mimicry. In *International Conference on Learning Representations (ICLR)*, 2024.